

Large Foundation Models and Reasoning

Pipi Hu

Beijing Institute of Mathematical Sciences and Applications

October 8, 2025

The question of whether machines can think is like asking whether submarines can swim.

Outline

- 1 Introduction to Reasoning in Large Foundation Models
- 2 Chain-of-Thought Reasoning
- 3 Mathematical Reasoning
- 4 Logical Reasoning
- 5 Advanced Reasoning Techniques
- 6 Reasoning Evaluation
- 7 Future Directions
- 8 Conclusion

What is Reasoning?

Reasoning is the process of drawing logical conclusions from given premises

- **Deductive reasoning:** From general to specific
- **Inductive reasoning:** From specific to general
- **Abductive reasoning:** Best explanation for observations
- **Common sense reasoning:** Everyday logical inference

Key challenge: How can large language models perform complex reasoning tasks?

Types of Reasoning in AI

- ① **Mathematical reasoning**: Solving equations, proofs
- ② **Logical reasoning**: Propositional and first-order logic
- ③ **Causal reasoning**: Understanding cause-effect relationships
- ④ **Commonsense reasoning**: Everyday knowledge and inference
- ⑤ **Multi-step reasoning**: Chain of thought processes

Question: Can foundation models truly reason or just pattern match?

Outline

- 1 Introduction to Reasoning in Large Foundation Models
- 2 Chain-of-Thought Reasoning**
- 3 Mathematical Reasoning
- 4 Logical Reasoning
- 5 Advanced Reasoning Techniques
- 6 Reasoning Evaluation
- 7 Future Directions
- 8 Conclusion

Chain-of-Thought (CoT) Prompting

Chain-of-Thought prompting encourages models to show their reasoning process

Input $\rightarrow$ Step 1 $\rightarrow$ Step 2 $\rightarrow \dots \rightarrow$ Answer (1)

Example:

- Problem: "Sarah has 12 apples. She gives 3 to Tom and 4 to Mary. How many does she have left?"
- CoT: "Sarah starts with 12 apples. She gives 3 to Tom: $12 - 3 = 9$. She gives 4 to Mary: $9 - 4 = 5$. Sarah has 5 apples left."

Mathematical Formulation of CoT

Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$ be a dataset where:

- x_i : Input problem
- y_i : Ground truth answer

CoT Training Objective:

$$\mathcal{L}_{CoT} = \mathbb{E}_{(x,y) \sim \mathcal{D}} [-\log p(y|x, \text{CoT})] \quad (2)$$

where CoT represents the intermediate reasoning steps.

Key insight: Decompose complex reasoning into simpler steps

Benefits of Chain-of-Thought

- ① **Transparency:** Makes reasoning process visible
- ② **Error localization:** Easier to identify where reasoning fails
- ③ **Improved accuracy:** Especially for multi-step problems
- ④ **Human interpretability:** Aligns with human problem-solving

Empirical evidence: CoT improves performance on:

- Mathematical word problems
- Logical reasoning tasks
- Commonsense reasoning
- Symbolic manipulation

Outline

- 1 Introduction to Reasoning in Large Foundation Models
- 2 Chain-of-Thought Reasoning
- 3 Mathematical Reasoning**
- 4 Logical Reasoning
- 5 Advanced Reasoning Techniques
- 6 Reasoning Evaluation
- 7 Future Directions
- 8 Conclusion

Mathematical Problem Solving

Large language models can solve mathematical problems through:

- ① **Step-by-step decomposition**
- ② **Algebraic manipulation**
- ③ **Logical inference**
- ④ **Pattern recognition**

Example - Quadratic Equation:

$$x^2 + 5x + 6 = 0 \quad (3)$$

$$(x + 2)(x + 3) = 0 \quad (4)$$

$$x = -2 \text{ or } x = -3 \quad (5)$$

Challenge: Ensuring mathematical correctness in generated solutions

Formal Mathematical Reasoning

For formal mathematical proofs, we need:

Given: $\Gamma \vdash \phi$ (premises) $\rightarrow$ Prove: ψ (conclusion) (6)

Proof strategies:

- Direct proof: $\Gamma \vdash \phi \rightarrow \psi$
- Proof by contradiction: $\Gamma, \neg\psi \vdash \perp$
- Proof by induction: Base case + inductive step
- Proof by cases: Exhaustive enumeration

Key challenge: Ensuring logical validity of each step

Mathematical Reasoning Capabilities

What LLMs can do:

- Solve arithmetic problems
- Perform algebraic manipulations
- Understand mathematical concepts
- Generate mathematical explanations

Current limitations:

- Inconsistent mathematical accuracy
- Difficulty with complex proofs
- Limited formal verification
- Hallucination of mathematical facts

Research direction: Improving mathematical reasoning reliability

Outline

- 1 Introduction to Reasoning in Large Foundation Models
- 2 Chain-of-Thought Reasoning
- 3 Mathematical Reasoning
- 4 Logical Reasoning**
- 5 Advanced Reasoning Techniques
- 6 Reasoning Evaluation
- 7 Future Directions
- 8 Conclusion

Propositional Logic Reasoning

Foundation models can perform logical inference using:

$$\text{Premises: } P_1, P_2, \dots, P_n \vdash \text{Conclusion: } Q \quad (7)$$

Logical connectives:

$$\text{Negation: } \neg P \quad (8)$$

$$\text{Conjunction: } P \wedge Q \quad (9)$$

$$\text{Disjunction: } P \vee Q \quad (10)$$

$$\text{Implication: } P \rightarrow Q \quad (11)$$

$$\text{Biconditional: } P \leftrightarrow Q \quad (12)$$

Example: If it rains, then the ground is wet. It's raining. Therefore, the ground is wet.

First-Order Logic Reasoning

For more complex reasoning involving quantifiers:

$$\text{Universal: } \forall x, P(x) \quad (13)$$

$$\text{Existential: } \exists x, P(x) \quad (14)$$

Logical inference rules:

- Modus Ponens: $P \rightarrow Q, P \vdash Q$
- Modus Tollens: $P \rightarrow Q, \neg Q \vdash \neg P$
- Universal Instantiation: $\forall x P(x) \vdash P(c)$
- Existential Generalization: $P(c) \vdash \exists x P(x)$

Challenge: Scaling logical reasoning to complex domains

Commonsense Reasoning

Everyday logical inference that humans perform naturally:

- **Causal reasoning:** Understanding cause-effect relationships
- **Temporal reasoning:** Understanding time sequences
- **Spatial reasoning:** Understanding spatial relationships
- **Social reasoning:** Understanding social interactions

Example: " John went to the store. He bought milk."

- Implicit: John had money, store was open, milk was available
- Inference: John now has milk, spent money, was at store

Key insight: Commonsense requires world knowledge and inference

Outline

- 1 Introduction to Reasoning in Large Foundation Models
- 2 Chain-of-Thought Reasoning
- 3 Mathematical Reasoning
- 4 Logical Reasoning
- 5 Advanced Reasoning Techniques**
- 6 Reasoning Evaluation
- 7 Future Directions
- 8 Conclusion

Tree of Thoughts (ToT)

Extends CoT by exploring multiple reasoning paths:

Problem $\rightarrow$ Multiple CoT paths $\rightarrow$ Best solution (15)

Process:

- ① Generate multiple reasoning paths
- ② Evaluate each path
- ③ Select the best path
- ④ Refine if necessary

Advantages:

- Explores multiple solution strategies
- Can backtrack from wrong paths
- More robust than single-path reasoning

Self-Consistency

Ensures reasoning consistency by:

Multiple reasoning attempts $\rightarrow$ Consensus answer

(16)

Method:

- ① Generate multiple reasoning paths for same problem
- ② Check consistency across paths
- ③ Select most consistent answer
- ④ Flag inconsistencies for review

Benefits:

- Reduces reasoning errors
- Identifies uncertain cases
- Improves overall reliability

Program-Aided Language Models (PAL)

Combines natural language reasoning with code execution:

Problem $\rightarrow$ Generate code $\rightarrow$ Execute $\rightarrow$ Answer (17)

Example:

- Problem: "Calculate the area of a circle with radius 5"
- Generated code: `import math; area = math.pi * 5**2`
- Execution: Returns 78.54
- Answer: "The area is approximately 78.54 square units"

Advantages:

- Precise mathematical computation
- Verifiable results
- Handles complex calculations

Outline

- 1 Introduction to Reasoning in Large Foundation Models
- 2 Chain-of-Thought Reasoning
- 3 Mathematical Reasoning
- 4 Logical Reasoning
- 5 Advanced Reasoning Techniques
- 6 Reasoning Evaluation**
- 7 Future Directions
- 8 Conclusion

Evaluating Reasoning Capabilities

Key metrics for reasoning evaluation:

- ① **Accuracy:** Correctness of final answers
- ② **Consistency:** Logical coherence of reasoning steps
- ③ **Completeness:** Coverage of all necessary steps
- ④ **Efficiency:** Optimal reasoning path length

Evaluation datasets:

- GSM8K: Grade school math problems
- MATH: High school mathematics
- LogiQA: Logical reasoning
- CommonsenseQA: Commonsense reasoning

Reasoning Error Analysis

Common reasoning errors:

- ① **Logical fallacies:** Invalid inference patterns
- ② **Factual errors:** Incorrect knowledge application
- ③ **Calculation mistakes:** Arithmetic errors
- ④ **Incomplete reasoning:** Missing steps

Error detection methods:

- Consistency checking across multiple attempts
- Step-by-step verification
- External knowledge validation
- Human expert review

Outline

- 1 Introduction to Reasoning in Large Foundation Models
- 2 Chain-of-Thought Reasoning
- 3 Mathematical Reasoning
- 4 Logical Reasoning
- 5 Advanced Reasoning Techniques
- 6 Reasoning Evaluation
- 7 Future Directions**
- 8 Conclusion

Challenges in Reasoning

Current limitations:

- ① **Hallucination:** Generating plausible but incorrect reasoning
- ② **Inconsistency:** Contradictory reasoning within same response
- ③ **Context limitations:** Difficulty with very long reasoning chains
- ④ **Domain specificity:** Poor generalization across domains

Research directions:

- Improving reasoning reliability
- Scaling to more complex problems
- Better evaluation methodologies
- Integration with external tools

Promising research directions:

- ① **Neuro-symbolic reasoning**: Combining neural and symbolic methods
- ② **Multi-agent reasoning**: Collaborative problem solving
- ③ **Interactive reasoning**: Human-AI reasoning collaboration
- ④ **Formal verification**: Mathematical proof of reasoning correctness

Key questions:

- Can we achieve truly reliable reasoning?
- How to scale reasoning to complex domains?
- What is the role of human oversight?

Outline

- 1 Introduction to Reasoning in Large Foundation Models
- 2 Chain-of-Thought Reasoning
- 3 Mathematical Reasoning
- 4 Logical Reasoning
- 5 Advanced Reasoning Techniques
- 6 Reasoning Evaluation
- 7 Future Directions
- 8 Conclusion**

Key insights about reasoning in foundation models:

- ① Chain-of-thought prompting significantly improves reasoning
- ② Mathematical and logical reasoning show promising results
- ③ Commonsense reasoning remains challenging
- ④ Evaluation and error detection are crucial

Future outlook:

- Continued improvement in reasoning capabilities
- Better integration with external tools
- More robust evaluation methods
- Potential for truly reliable AI reasoning

Research questions for the community:

- ① How can we ensure reasoning reliability at scale?
- ② What are the fundamental limits of neural reasoning?
- ③ How to best combine symbolic and neural approaches?
- ④ What role should human oversight play?

Contact: For questions and collaboration, please reach out at (hpp@bimsa.cn)

Thank you

Questions and Discussion

Thank you for your attention!